

✅ 1. Problem Scope (5 points)

- **Problem:** Hospital readmissions within 30 days indicate potential care gaps, increase costs, and affect patient outcomes. Predicting patients at high risk of readmission helps improve post-discharge care.
- **Objective:**
Build an AI model that predicts the likelihood of a patient being readmitted within 30 days of discharge. The model will assist clinicians in prioritizing follow-ups, reducing readmission rates, and optimizing resource allocation.
- **Stakeholders:**
 - **Hospital administrators** – interested in lowering costs and improving KPIs
 - **Clinicians** – want actionable insights to improve care
 - **Patients** – benefit from proactive interventions
 - **Data privacy/compliance officers** – ensure the system adheres to regulations

✅ 2. Data Strategy (10 points)

📊 a) Data Sources:

- **Electronic Health Records (EHR):** diagnoses, vitals, treatments, discharge summaries
- **Demographics:** age, gender, socioeconomic status, residence
- **Lab Results:** blood tests, imaging results
- **Medication History:** prescriptions, dosage, adherence
- **Past Admission Records:** frequency and reasons for previous visits
- **Doctor and Nurse Notes:** extracted via NLP (optional advanced step)

⚖️ b) Ethical Concerns:

1. **Patient Privacy & Confidentiality**
 - EHRs contain sensitive personal information. Data must be anonymized, encrypted, and handled under **HIPAA** or local privacy laws.
2. **Bias in Training Data**

- If historical records reflect biased care (e.g., lower quality for underserved groups), the model could **unfairly predict higher readmission risks** for those patients. This must be mitigated through fairness-aware learning.

1. Data Collection

2. Data Cleaning:

- Remove duplicates
- Handle missing values (e.g., mean imputation for vitals)

3. Feature Engineering:

- Time since last visit
- Number of chronic conditions
- Count of emergency visits in past year
- Age group buckets (e.g., 18-40, 41-65, 66+)

4. Encoding:

- One-hot encode categorical variables (e.g., discharge type)

5. Normalization:

- Scale lab results and vitals

6. Label Definition:

- Readmitted within 30 days → `1`
- Not readmitted → `0`

7. Split dataset: 70% train / 30% test

1. Data Collection

2. Data Cleaning:

- Remove duplicates
- Handle missing values (e.g., mean imputation for vitals)

3. Feature Engineering:

- Time since last visit
- Number of chronic conditions

- Count of emergency visits in past year
- Age group buckets (e.g., 18-40, 41-65, 66+)

4. Encoding:

- One-hot encode categorical variables (e.g., discharge type)

5. Normalization:

- Scale lab results and vitals

6. Label Definition:

- Readmitted within 30 days → `1`
- Not readmitted → `0`

7. Split dataset: 70% train / 30% test

3. Model Development (10 points)

Model Selection:

Random Forest Classifier

- **Why:**
 - Handles both numerical and categorical data
 - Naturally handles missing data and outliers
 - Provides feature importance (explainability)
 - Performs well on imbalanced datasets (with class weighting)

Hypothetical Confusion Matrix:

	Predicted Readmit	Predicted No Readmit
Actual Readmit	80 (TP)	20 (FN)
Actual No Readmit	30 (FP)	170 (TN)

Evaluation:

- **Precision** = $TP / (TP + FP) = 80 / (80 + 30) = \mathbf{0.727}$

- **Recall** = $TP / (TP + FN) = 80 / (80 + 20) = \mathbf{0.800}$
- **F1 Score** = $2 * (Precision * Recall) / (Precision + Recall) = \mathbf{0.761}$

Readmit - (short for *readmission*) refers to a patient being re-admitted to the hospital within 30 days after being discharged.

✅ 4. Deployment (10 points)

🔑 Integration Steps:

1. **Model API:** Wrap the trained model as a REST API using Flask or FastAPI.
2. **Authentication:** Secure the API using OAuth2 or hospital SSO systems.
3. **EHR System Integration:**
 - Integrate the API with the hospital's EHR interface (e.g., Epic, Cerner)
 - Show readmission risk score on the patient discharge screen
4. **Feedback Loop:**
 - Allow doctors to flag false positives/negatives for future model retraining.
5. **Monitoring:** Log model predictions and performance for audit and drift detection.

⚖️ Regulatory Compliance:

- **HIPAA Alignment:**
 - Ensure data encryption (at rest and in transit)
 - Log data access events
 - Store models and predictions securely
 - Conduct privacy impact assessments before launch

✅ 5. Optimization (5 points)

🚫 How to Reduce Overfitting:

Use Cross-Validation + Regularization

- Apply **k-fold cross-validation** during training to ensure the model performs consistently across all data segments.

- Add **regularization** (e.g., limit tree depth or use dropout in neural nets).
- Bonus: Perform **feature selection** to eliminate noisy inputs that confuse the model.

This hospital AI project reflects how data-driven precision can transform patient care. By aligning clinical goals with AI best practices—while staying compliant and ethical—we can deliver a system that empowers decision-makers, saves lives, and adapts to real-world challenges. With intelligent preprocessing, transparent model choices, and seamless deployment, this solution embodies both **technical depth** and **human-centered design** — the heart of what Tech Finesse stands for.

Part 3: Critical Thinking

Ethics & Bias (10 points)

How might biased training data affect patient outcomes?

If the AI model is trained on **biased historical data**, it can reinforce or even amplify existing healthcare disparities. For example:

- **Underrepresented groups** (e.g., rural patients, certain ethnicities) may have received fewer diagnostic tests or follow-ups in the past.
- The model could **learn that pattern as “normal”**, and then:
 - Predict **lower readmission risk** for those groups (when in fact, risk is high)
 - Deny timely follow-up care → leading to worse health outcomes

This results in **algorithmic discrimination**, where the AI favors certain populations over others, potentially **endangering vulnerable patients**.

Strategy to Mitigate Bias:

Use Fairness-Aware Data Preprocessing

- Analyze and balance the training data by:
 - **Stratifying by gender, ethnicity, location**
 - **Up-sampling underrepresented groups** or applying **re-weighting techniques**
- Apply **bias detection metrics** (e.g., disparate impact ratio) during training

- Involve **clinical experts** from diverse backgrounds to audit predictions

Think of it as embedding “ethical QA” into the model pipeline.

Trade-offs (10 points)

1. Model Interpretability vs. Accuracy in Healthcare

Definition of the Trade-off

In machine learning, there is often a tension between **interpretability** and **accuracy**:

- **Interpretability** means how easily humans (especially non-technical users like clinicians) can understand how a model arrives at its predictions.
- **Accuracy** reflects how well the model performs (e.g., predicts patient readmission correctly).

In many domains, a highly accurate "black-box" model may be acceptable. But **in healthcare**, decisions must be:

- **Transparent**
 - **Ethically sound**
 - **Legally defensible**
-

Why This Trade-off Matters in Healthcare

1. Regulatory Requirements & Accountability

- Healthcare is governed by strict regulations (e.g., HIPAA, GDPR).
- Doctors and hospitals must be able to **justify decisions**, especially when those decisions affect patient outcomes.
- Using a model that can't be explained (like a deep neural net) may **violate compliance or expose institutions to legal risk**.

2. Trust and Adoption by Clinicians

- Doctors are more likely to **trust and use** a model they understand.
 - Example: A logistic regression model saying, “The patient is at 80% risk because of recent ER visits and chronic kidney disease” is more acceptable than a black box that outputs “80%” with no rationale.

- Low interpretability leads to **resistance in adoption**, even if accuracy is high.

3. Clinical Safety and Risk Management

- **False positives** could lead to unnecessary tests or hospitalizations.
- **False negatives** could result in missing critical interventions.
- Interpretable models help healthcare providers **spot and mitigate such risks** before acting on them.
- **High Accuracy Models** (e.g., deep neural networks, ensembles):
 - Great at prediction, but difficult to interpret (“black box”)
 - Problem: Doctors can’t explain why the AI flagged a patient → **trust gap**
- **High Interpretability Models** (e.g., decision trees, logistic regression):
 - Easier to understand → clinicians can verify and trust predictions
 - But may **sacrifice accuracy** in complex cases

In healthcare, **explainability often outweighs raw accuracy**, because decisions affect human lives and must be accountable.

2. If the Hospital Has Limited Computational Resources...

1. Prefer Lightweight, Interpretable Models

 **Use models like Logistic Regression, Decision Trees, or Naive Bayes.**

These models: Require minimal processing power, train quickly even on basic machine and are easier to audit and explain — which is vital in healthcare

They may not capture all deep patterns, but they’re **efficient, practical, and safe for deployment** on limited infrastructure.

2. Avoid Deep Learning Unless Offloaded to the Cloud

 **Neural networks and ensemble models (e.g., XGBoost) are powerful but resource-heavy.**

Running them locally requires: High-performance CPUs/GPUs, more RAM and longer training and inference times.

Solution: If deep learning is essential, deploy it via, cloud services (e.g., AWS SageMaker, Google AI Platform) and Lightweight inference tools (e.g., TensorFlow Lite, ONNX Runtime). This shifts the computational load **off-site**, keeping local systems light.

3. Use Batch Prediction Instead of Real-Time Processing

 **Batch processing = Efficient use of resources**

Instead of real-time predictions (which demand constant availability): Run predictions **at scheduled times** (e.g., every 4 hours or nightly), queue new patient data, process in bulk. This reduces server strain and power consumption

4. Optimize the Feature Set (Less is More)

 **Reduce the number of input variables (features) to streamline computation.**

Focus only on **top contributing features** (e.g., length of stay, prior admissions, chronic conditions)

5. Deploy Using Efficient Toolchains and Formats

 **Choose tools that are optimized for low-resource environments.**

- Use **Flask** or **FastAPI** to deploy lightweight ML services, store models in **compressed formats** (e.g., .joblib, .pkl, .onnx), run on edge-friendly environments (e.g., Raspberry Pi, hospital internal servers)

Also ensure: Minimal dependencies, efficient logging, caching mechanisms for repeat queries. This ensures **cost-effective, reliable model serving** without overloading the system.

Part 4: Reflection & Workflow Diagram (10 points)

 **Reflection (5 points)**

What was the most challenging part of the workflow? Why?

The most challenging part of the workflow was the **data preprocessing and feature engineering** phase. Cleaning clinical data requires domain-specific understanding, careful handling of missing or inconsistent records, and proper transformation of complex health metrics.

Additionally, ensuring that the dataset remained **bias-free and privacy-compliant** added another layer of complexity. Without real patient data, creating meaningful **hypothetical datasets** that simulate realistic trends was intellectually demanding and time-consuming.

How would you improve your approach with more time/resources?

With more time and resources, I would:

- Collaborate with a medical expert or data steward to validate feature relevance and clinical meaning.
- Use advanced **automated feature engineering tools** (like Featuretools or DataRobot) to generate richer features.
- Integrate real-world anonymized healthcare datasets (e.g., from MIMIC-III or open EHR repositories).
- Deploy the model via a secure cloud-based dashboard with real-time risk alerts for clinicians — closing the loop from prediction to action.

This would make the workflow more robust, ethical, and production-ready.